

DMLab2 report

113062571 CS 張原鳴



Data Preparation

Initial data: Each row data corresponds to one tweet id

- data identification.csv : Classify the data into train and test data
- emotion.csv : Label the data in different kind of emotion
- tweets_DM.json : Include _score, _index, _source, _crawldata and _type, _source include:
 - tweet_id
 - hashtag
 - text

Data Preparation

I extract the csv files into train_data.csv and test_data.csv in data_create.py.

- Feature selection:
 - id
 - text
 - hashtag
 - _score
- I focus on only text part during the training in the end.

```
import pandas as pd
df = pd.read_json('tweets_DM.json', lines=True)
#print(df)
data_identification = pd.read_csv('data_identification.csv')
data_emotion = pd.read_csv('emotion.csv')
#print(data_identification['identification'].value_counts())
#print(data_identification)
src = df['_source']
df_src = pd.DataFrame([
    'hashtags': item['tweet']['hashtags'],
    'tweet_id': item['tweet']['tweet_id'],
    'text': item['tweet']['text']
] for item in src)
tweet_df = pd.concat([df_src, df['_score']], axis=1)
#print(df_merged)
#print(df_merged['identification'].value_counts())
df_result = pd.merge(tweet_df, data_identification, on='tweet_id')
#print(df_result)
df_train = df_result[df_result['identification'] == 'train']
df_test = df_result[df_result['identification'] == 'test']
df_train = pd.merge(df_train, data_emotion, on='tweet_id')

df_train = df_train[['tweet_id', 'text', 'hashtags', '_score', 'emotion']]
df_test = df_test[['tweet_id', 'text', 'hashtags', '_score']]

df_train.to_csv('train_data.csv', index=False)
df_test.to_csv('test_data.csv', index=False)

print(df_train)
print(df_test)
```

Data Preprocessing

- TfidfVectorizer:
 - Transform a collection of text data in train data into a numerical matrix based on the TF-IDF (Term Frequency-Inverse Document Frequency) score.
 - Limits the vocabulary size to the top 25,000 most important words across all documents, ranked by their TF-IDF score.
- Stop_word = 'english':
 - Automatically removes common English stop words (e.g., "the", "is", "and") from the text before computing TF-IDF, ignoring words that are meaningless.

```
In [5]: from sklearn.feature_extraction.text import TfidfVectorizer
        TFIDF = TfidfVectorizer(max_features=25000, stop_words='english')
        TFIDF.fit(train_df['text'])
```

```
Out[5]:
```

▼ TfidfVectorizer

TfidfVectorizer(max_features=25000, stop_words='english')

Data Preprocessing

- Split and transform the training data into train and validation part (8:2).
 - We can observe the results not only in training phase but also in validation phase.
- LabelEncoder:
 - Models cannot process text labels (e.g., “anticipation”) directly and need numerical inputs, so I use LabelEncoder for one-hot encoding.
 - Dimension of y_train: 1 to 8 (number of target class).

```
In [7]: X_train, X_val, y_train, y_val = train_test_split(
        X_train, y_train, test_size=0.2, random_state=42
        )
```

```
In [8]: label_encoder = LabelEncoder()
        label_encoder.fit(y_train)
```

```
Out[8]:
```

▼ LabelEncoder

LabelEncoder()

Create the Model

- Model: Deep neural network (DNN) with six hidden layers
- Regularization:
 - Batch Normalization ensures stable and faster training.
 - Dropout reduces overfitting and prevents the network from becoming overly reliant on specific neurons.

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 25000)	0
dense (Dense)	(None, 1024)	25,601,024
batch_normalization (BatchNormalization)	(None, 1024)	4,096
re_lu (ReLU)	(None, 1024)	0
dropout (Dropout)	(None, 1024)	0
dense_1 (Dense)	(None, 512)	524,800
batch_normalization_1 (BatchNormalization)	(None, 512)	2,048
re_lu_1 (ReLU)	(None, 512)	0
dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 256)	131,328
batch_normalization_2 (BatchNormalization)	(None, 256)	1,024
re_lu_2 (ReLU)	(None, 256)	0
dropout_2 (Dropout)	(None, 256)	0
dense_3 (Dense)	(None, 128)	32,896
batch_normalization_3 (BatchNormalization)	(None, 128)	512
re_lu_3 (ReLU)	(None, 128)	0
dense_4 (Dense)	(None, 64)	8,256
batch_normalization_4 (BatchNormalization)	(None, 64)	256
re_lu_4 (ReLU)	(None, 64)	0
dense_5 (Dense)	(None, 32)	2,080
batch_normalization_5 (BatchNormalization)	(None, 32)	128
re_lu_5 (ReLU)	(None, 32)	0
dense_6 (Dense)	(None, 8)	264

Training the Model

- Epoch: 30, batch size: 256
- Early_Stopping: Stop the training process when the model's performance on a validation dataset stops improving.
- Generate the output result into submission.csv

```
epochs = 30
batch_size = 256

from tensorflow.keras.callbacks import EarlyStopping

# Define EarlyStopping
early_stopping = EarlyStopping(
    monitor='val_accuracy',
    patience=2,
    mode='max',
    restore_best_weights=True
)

# training!
history = model.fit(
    X_train, y_train,
    epochs=epochs,
    batch_size=batch_size,
    validation_data=(X_val, y_val),
    callbacks=[early_stopping]
)
```

Trying for Improvement

I try some methods, but the score just dropped, so I put in here.

- Sampling partial data based on frac
- Process data balancing.

Counts of every class in training data:

Speculation of the worse result:

1. Bad quality of generating data
because text data is hard to simulate.
2. Overfitting problem.

emotion	
joy	516017
anticipation	248935
trust	205478
sadness	193437
disgust	139101
fear	63999
surprise	48729
anger	39867

Name: count, dtype: int64



emotion	
surprise	516017
disgust	516017
trust	516017
anger	516017
fear	516017
sadness	516017
anticipation	516017
joy	516017

Name: count, dtype: int64

Trying for Improvement

- Try different models:
 - Random Forest
 - Transformer embeddings
- Speculation of the worse result:
 - The size of training data is too big, Random Forests are effective on smaller datasets, they may not scale well with such a large amount of data.
 - Lack of fine-tune in pretrained Transformer model, may not align perfectly with our task.

Thanks